



Top LLM Providers 2026

Ranked by capability · Best models · Pricing per 1M tokens

● #1 Anthropic — Quality Leader

- Best model: Claude Opus 4.7 · Best for agents, long reasoning, coding
 - Opus 4.7: \$5 / \$25 · Sonnet 4.6: \$3 / \$15 · Haiku 4.5: \$0.80 / \$4
 - 200K context · 90% cache discount · 7-hour agentic sessions
-

● #2 OpenAI — Broadest Portfolio

- Best model: GPT-5.5 · Best for tool use, multi-step planning, agents
 - GPT-5.5: \$10 / \$30 · GPT-5.4: \$2.50 / \$15 · Nano: \$0.10 / \$0.40
 - 150× price range within one provider · 88.7% SWE-bench (Codex)
-

● #3 Google — Best Free Tier

- Best model: Gemini 3 Pro · 2M token context window leads all providers
 - Gemini 3 Pro: \$2 / \$12 · Flash: \$0.30 / \$2.50 · Flash-Lite: free tier
 - 1,000 free API calls/day · Best value mid-tier in the market
-

● #4 xAI — Fast Riser

- Best model: Grok 4.3 · Trained on 200K-GPU Colossus cluster
- Grok 4.3: \$3 / \$15 · Grok Mini: \$0.20 / \$0.50 · API pricing competitive
- Strong reasoning · Included in X Premium+ subscription

● #5 DeepSeek — Best Value

- Best model: DeepSeek V3.2 · Rivals models 10× its price on benchmarks
- V3.2: \$0.14 / \$0.28 · R1 reasoning free on OpenRouter
- Open-weights · Self-hostable · Cheapest high-quality API available

● #6 Mistral — Best European Option

- Best model: Mistral Large 3 · GDPR-native, EU data residency
- Large 3: ~\$2 / \$6 · Mistral Small 4: budget tier · Codestral: code-only
- OpenAI-compatible API · Strong on multilingual & enterprise compliance

● #7 Meta Llama — Open Source King

- Best model: Llama 4 Maverick · Free weights, self-host or use via API
- Via Together AI / Groq: ~\$0.18–\$0.90 per 1M tokens
- No per-token fees self-hosted · Breaks even vs APIs at 500K+ tok/day

● ☐ Price-to-Performance Sweet Spots

- Budget: DeepSeek V3.2 (\$0.14) or Gemini Flash (\$0.30) — best per-dollar
- Balanced: Claude Sonnet 4.6 (\$3) or GPT-5.4 (\$2.50) — prod standard
- Power: Claude Opus 4.7 or GPT-5.5 — long agents & hard reasoning

All prices per 1M tokens (input / output). Cached inputs save up to 90%.
Routing cheap models for simple tasks cuts costs 50–80%.